

Consider a dataset $D = (\mathbf{x}_i, y_i)^N$ of N data points, where $\mathbf{x}_i = (x_{i1}, x_{i2}, \dots, x_{iM})$ is a **feature vector** with M features, and y_i is the **target**, i.e., the response, variable. Let x_j denote the j th variable in feature space. A **typical linear regression model** can then be expressed mathematically as:

$$\begin{array}{ccccccc}
 & & \text{slope} & & \text{feature} & & \\
 & & \downarrow & & \downarrow & & \\
 y = & \beta_0 & + & \beta_1 x_1 & + & \beta_2 x_2 & + \dots + \beta_M x_M \\
 \uparrow & & & & & & \\
 \text{target} & & & & & &
 \end{array}$$

This model assumes that the relationships between the target variable y_i and features x_j are linear and can be captured in slope terms $\beta_1, \beta_2, \dots, \beta_M$.